

```
from google.colab import files
uploaded = files.upload()
```

Choose Files | 4 files

X_test.csv(text/csv) - 334426 bytes, last modified: 4/4/2026 - 100% done
X_train.csv(text/csv) - 1352947 bytes, last modified: 4/4/2026 - 100% done
y_test.csv(text/csv) - 16040 bytes, last modified: 4/4/2026 - 100% done
y_train.csv(text/csv) - 64115 bytes, last modified: 4/4/2026 - 100% done
Saving X_test.csv to X_test.csv
Saving X_train.csv to X_train.csv
Saving y_test.csv to y_test.csv
Saving y_train.csv to y_train.csv

```
!pip install -q transformers datasets accelerate scikit-learn
```

```
import pandas as pd
import numpy as np
import torch

from transformers import (
    DistilBertTokenizerFast,
    DistilBertForSequenceClassification,
    Trainer,
    TrainingArguments
)
from sklearn.metrics import accuracy_score, precision_recall_fscore_support
```

```
X_train = pd.read_csv("X_train.csv")
X_test = pd.read_csv("X_test.csv")
y_train = pd.read_csv("y_train.csv")
y_test = pd.read_csv("y_test.csv")

print(X_train.head())
print(y_train.head())
print(X_train.shape, X_test.shape, y_train.shape, y_test.shape)
```

```

                                headline
0                                fbi raids fridge
1      old el paso introduces emergency taco kit
2  larry kudlow 'leaning' toward senate run in co...
3  'everything's $10,000' chain goes out of business
4  what this ceo did proves that introverts make ...
   is_sarcastic
0              1
1              1
2              0
3              1
4              0
(21367, 1) (5342, 1) (21367, 1) (5342, 1)
```

```
train_texts = X_train["headline"].tolist()
test_texts = X_test["headline"].tolist()

train_labels = y_train["is_sarcastic"].tolist()
test_labels = y_test["is_sarcastic"].tolist()

print(train_texts[:3])
print(train_labels[:3])
```

```
['fbi raids fridge', 'old el paso introduces emergency taco kit', 'larry kudlow 'leaning' toward senate run in connecticut']
[1, 1, 0]
```

```
print("CUDA available:", torch.cuda.is_available())
if torch.cuda.is_available():
    print("GPU:", torch.cuda.get_device_name(0))
```

```
CUDA available: True
GPU: Tesla T4
```

```
tokenizer = DistilBertTokenizerFast.from_pretrained("distilbert-base-uncased")
```

```

/usr/local/lib/python3.12/dist-packages/huggingface_hub/utils/_auth.py:94: UserWarning:
The secret `HF_TOKEN` does not exist in your Colab secrets.
To authenticate with the Hugging Face Hub, create a token in your settings tab (https://huggingface.co/settings/tokens), set it
You will be able to reuse this secret in all of your notebooks.
Please note that authentication is recommended but still optional to access public models or datasets.
warnings.warn(

tokenizer_config.json: 100%                               48.0/48.0 [00:00<00:00, 2.50kB/s]

Warning: You are sending unauthenticated requests to the HF Hub. Please set a HF_TOKEN to enable higher rate limits and faster d
WARNING:huggingface_hub.utils._http:Warning: You are sending unauthenticated requests to the HF Hub. Please set a HF_TOKEN to en
vocab.txt:      232k/? [00:00<00:00, 3.02MB/s]

tokenizer.json:    466k/? [00:00<00:00, 7.75MB/s]

```

```

train_encodings = tokenizer(
    train_texts,
    truncation=True,
    padding=True,
    max_length=64
)

test_encodings = tokenizer(
    test_texts,
    truncation=True,
    padding=True,
    max_length=64
)

```

```

class SarcasmDataset(torch.utils.data.Dataset):
    def __init__(self, encodings, labels):
        self.encodings = encodings
        self.labels = labels

    def __getitem__(self, idx):
        item = {key: torch.tensor(val[idx]) for key, val in self.encodings.items()}
        item["labels"] = torch.tensor(self.labels[idx], dtype=torch.long)
        return item

    def __len__(self):
        return len(self.labels)

```

```

train_dataset = SarcasmDataset(train_encodings, train_labels)
test_dataset = SarcasmDataset(test_encodings, test_labels)

```

```

model = DistilBertForSequenceClassification.from_pretrained(
    "distilbert-base-uncased",
    num_labels=2
)

```

```

config.json: 100%                               483/483 [00:00<00:00, 38.9kB/s]

model.safetensors: 100%                           268M/268M [00:02<00:00, 244MB/s]

Loading weights: 100%                             100/100 [00:00<00:00, 473.98it/s, Materializing param=distilbert.transformer.layer.5.sa_layer_norm.weight]
DistilBertForSequenceClassification LOAD REPORT from: distilbert-base-uncased

```

Key	Status
vocab_projector.bias	UNEXPECTED
vocab_transform.bias	UNEXPECTED
vocab_layer_norm.weight	UNEXPECTED
vocab_layer_norm.bias	UNEXPECTED
vocab_transform.weight	UNEXPECTED
pre_classifier.weight	MISSING
classifier.bias	MISSING
pre_classifier.bias	MISSING
classifier.weight	MISSING

```

Notes:
- UNEXPECTED :can be ignored when loading from different task/architecture; not ok if you expect identical arch.
- MISSING :those params were newly initialized because missing from the checkpoint. Consider training on your downstream task.

```

```

def compute_metrics(eval_pred):
    logits, labels = eval_pred

```

```

preds = np.argmax(logits, axis=1)

precision, recall, f1, _ = precision_recall_fscore_support(
    labels, preds, average="binary"
)
acc = accuracy_score(labels, preds)

return {
    "accuracy": acc,
    "precision": precision,
    "recall": recall,
    "f1": f1
}

```

```

training_args = TrainingArguments(
    output_dir="./results",
    eval_strategy="epoch",
    save_strategy="epoch",
    logging_strategy="epoch",
    num_train_epochs=3,
    per_device_train_batch_size=16,
    per_device_eval_batch_size=16,
    weight_decay=0.01,
    load_best_model_at_end=True,
    metric_for_best_model="f1",
    greater_is_better=True,
    report_to="none"
)

```

```

trainer = Trainer(
    model=model,
    args=training_args,
    train_dataset=train_dataset,
    eval_dataset=test_dataset,
    compute_metrics=compute_metrics
)

```

```
trainer.train()
```

 [4008/4008 08:04, Epoch 3/3]

Epoch	Training Loss	Validation Loss	Accuracy	Precision	Recall	F1
1	0.274673	0.196359	0.926432	0.903306	0.932196	0.917524
2	0.105100	0.278491	0.929614	0.899716	0.944989	0.921797
3	0.029755	0.351363	0.934669	0.939648	0.909595	0.924377

Writing model shards: 100% 1/1 [00:03<00:00, 3.11s/it]

Writing model shards: 100% 1/1 [00:05<00:00, 5.91s/it]

Writing model shards: 100% 1/1 [00:01<00:00, 1.14s/it]

There were missing keys in the checkpoint model loaded: ['distilbert.embeddings.LayerNorm.weight', 'distilbert.embeddings.LayerN

There were unexpected keys in the checkpoint model loaded: ['distilbert.embeddings.LayerNorm.beta', 'distilbert.embeddings.Layer

TrainOutput(global_step=4008, training_loss=0.13650926144537098, metrics={'train_runtime': 486.3593, 'train_samples_per_second': 131.798, 'train_steps_per_second': 8.241, 'total_flos': 1061411590156032.0, 'train_loss':

```

eval_results = trainer.evaluate()
eval_results

```

 [334/334 00:07]

```

{'eval_loss': 0.351362943649292,
 'eval_accuracy': 0.9346686634219393,
 'eval_precision': 0.939647577092511,
 'eval_recall': 0.9095948827292111,
 'eval_f1': 0.924377031419285,
 'eval_runtime': 7.6897,
 'eval_samples_per_second': 694.697,
 'eval_steps_per_second': 43.435,
 'epoch': 3.0}

```

```

pred_output = trainer.predict(test_dataset)
preds = np.argmax(pred_output.predictions, axis=1)

pred_df = pd.DataFrame({

```

```
"headline": test_texts,
"actual": test_labels,
"predicted": preds
})

pred_df.head()
```

	headline	actual	predicted
0	the outsider has officially squeezed its way i...	0	0
1	lgbtq activists organizing massive dance prote...	0	0
2	romney dominated debate, say pundits trying to...	1	1
3	israeli soldiers open fire on palestinians car...	1	0
4	learned sage points out that powerball not as ...	1	1

Next steps:

[Generate code with pred_df](#)

[New interactive sheet](#)

```
misclassified = pred_df[pred_df["actual"] != pred_df["predicted"]]
misclassified.head(10)
```

	headline	actual	predicted
3	israeli soldiers open fire on palestinians car...	1	0
7	americans finally found something to drink tha...	0	1
8	that show about the lady sheriff finally relea...	1	0
51	defense department typo results in u.s. attack...	1	0
66	'fire rainbow' supplants double rainbow as soc...	0	1
103	former presidents convene for liver spot summit	1	0
105	twitter creator on iran: 'i never intended for...	1	0
131	u.s. military prepares for biggest okinawa lan...	0	1
143	u.s. improves infrastructure with transnationa...	1	0
161	police finally make breakthrough in decades-ol...	1	0

Next steps:

[Generate code with misclassified](#)

[New interactive sheet](#)

```
correct = pred_df[pred_df["actual"] == pred_df["predicted"]]
correct.head(10)
```

	headline	actual	predicted
0	the outsider has officially squeezed its way i...	0	0
1	lgbtq activists organizing massive dance prote...	0	0
2	romney dominated debate, say pundits trying to...	1	1
4	learned sage points out that powerball not as ...	1	1
5	all about otters!	0	0
6	mitt romney jots down ideas for concession spe...	1	1
9	shamed and abandoned: the fate of syria's form...	0	0
10	women leaders talk personal: how to be a true ...	0	0
11	survival myths that could actually kill you	0	0
12	letter from employer thankfully omits balls-co...	1	1

Next steps:

[Generate code with correct](#)

[New interactive sheet](#)

```
pred_df.to_csv("bert_predictions.csv", index=False)
```

```
metrics_df = pd.DataFrame([eval_results])
metrics_df.to_csv("bert_eval_results.csv", index=False)
```

```
from google.colab import files

files.download("bert_predictions.csv")
files.download("bert_eval_results.csv")
```

Start coding or [generate](#) with AI.